

Abdul Moiz

AI / ML Engineer

+92-319-8138264 | Karachi, Pakistan | abdulmoiz.aiml.dev@gmail.com

[Portfolio](#) | [LinkedIn](#) | [GitHub](#)

Profile

Final-year Software Engineering student building end-to-end AI systems: explainable loan-risk prediction, fine-tuned coding assistance, and medical evidence retrieval. Experience spans model evaluation, inference APIs, testing, and Docker packaging.

Education

University of Karachi (UBIT) | Karachi, Pakistan

Bachelor in Software Engineering | Expected Dec 2027 | CGPA: **3.30 / 4.00**

Selected Projects

Loan Risk Assessment System

Feb - Mar 2026

Explainable credit-risk prediction | [GitHub](#) | [Live Demo](#)

- Built a loan-risk application that predicts default risk, estimates interest rates, and segments borrowers; translated SHAP explanations into plain English to make predictions easier to interpret.
- Trained on the Kaggle Credit Risk Dataset (32,581 records). LightGBM achieved **0.947 ROC AUC** and **0.839 F1** on **2,444 test records**, versus 0.873 AUC for logistic regression.
- Deployed a Docker-packaged FastAPI/Streamlit application on Hugging Face Spaces, with ONNX model inference and PostgreSQL records for inputs, predictions, and explanations.

CodeMentor-LLM

Apr - Jun 2026

Fine-tuned coding assistant | [GitHub](#) | [Live Demo](#)

- Fine-tuned **Llama 3.2 3B** for coding assistance with QLoRA supervised training on **5,000 coding pairs**, followed by DPO alignment on 200 preference pairs.
- Improved **ROUGE-L** by **14.9%** versus the base model (0.1807 to 0.2077) in a **50-example test evaluation**; achieved BERTScore F1 of 0.8491 and tracked training in Weights & Biases.
- Deployed FastAPI inference on a Modal A10G GPU with a Hugging Face Spaces interface; added GitHub Actions CI/CD and **29 pytest checks** for data preparation, input validation, and API behavior.

MedRAG

Jul - Sep 2026

Medical evidence assistant / Local Docker application | [GitHub](#)

- Built a multilingual medical assistant with source citations, retrieving from **22,696 indexed chunks** across PubMed, WHO guidelines, and OpenFDA.
- Combined hybrid retrieval, reranking, and Neo4j graph enrichment; reduced pipeline latency by **49.8-66.5%** on three warmed local queries using ONNX, preserving top-five retrieval order.
- Added sign-in, saved chats, and chat-scoped PDF uploads; packaged FastAPI, Streamlit, and databases with Docker and tested authentication, isolation, and persistence.

Technical Skills

Programming & data: Python, SQL, Pandas, NumPy, scikit-learn

ML & LLMs: LightGBM, SHAP, PyTorch, Transformers, QLoRA, SFT, DPO

Retrieval & storage: Qdrant, BM25, Neo4j, PostgreSQL

Application delivery: FastAPI, Streamlit, Docker, ONNX Runtime, GitHub Actions, pytest